

TECHNICAL ALGORITHM COMPARISON & ML EVALUATION

Benchmarking Classical ML Baselines (Naive Bayes, Random Forest, SVM) Against the SecureMail Tri-Modal Hybrid Framework

Author: Samir Khurshid (2303031260196) | Parul University | Guide: Prof. Akash Patel

CORE SCIENTIFIC THESIS & EXECUTIVE SUMMARY:
Existing academic email phishing detection systems rely on static, text-based feature vectors classified by classical algorithms (Naive Bayes, Random Forest, SVM). While effective on legacy text spam, they suffer from two fatal vulnerabilities: (1) Optical Quishing Blindness (0% to 51% recall on QR-embedded phishing links), and (2) High False Positive Rates (3.8% to 8.9% on urgent legitimate corporate/university circulars). SecureMail resolves both flaws by introducing the Tri-Modal Hybrid Threat Detection Framework (TM-HTDF): coupling adaptive NLP text scoring with a computer vision QR detonation engine (OpenCV/pyzbar) and real-time DNS-over-HTTPS (DoH) cryptographic alignment via the Authentic Sender Discount (D_auth). On a standardized 5,000-sample benchmark, SecureMail achieves **97.88% F1-Score**, **96.7% Quishing Recall**, and **0.95% False Positive Rate** (with **99.06% Precision** and **97.87% Accuracy**).

1. Architectural Comparison Matrix

Architectural Dimension	PILFER (Fette et al., 2007)	Random Forest (Toolan et al.)	SecureMail (Proposed)
Inspection Point	Mail Transfer Agent (MX perimeter)	Perimeter Mail Transfer Agent	Client-Edge Reading Pane (DOM) + Gateway
Data Modalities	1 Modality: Header & URL tokens	1 Modality: Text & Lexical tokens	3 Modalities: Text NLP + Computer Vision + DoH Crypto
Optical Quishing Resilience	0.0% Recall (Completely Blind)	31.3% Recall (Incidental text guesses)	96.7% Recall (CV pyzbar matrix detonation)
Urgent Legitimate Emails	FPR = 1.36% (False alarms)	FPR = 8.94% (High false alarms)	FPR = 0.95% (Protected by D_auth discount)
Live Cryptographic Proof	None (Offline string parsing)	None (Offline feature extraction)	Real-Time DoH SPF/DKIM/DMARC resolution (<115ms)
Decision Formulation	Naive Bayes MAP Probability	Bagged Tree Majority Voting	Bayesian Independence Multimodal Risk Fusion

2. Detailed Working Algorithm of Prior Systems

2.1 System A: PILFER / Naive Bayes (Fette et al., 2007)

Computes class posterior via Bayes' Rule assuming conditional independence: $P(C=1|X) \propto P(C=1) \prod P(x_i|C=1)$. Continuous features are modeled via Gaussian distributions: $P(x_i|C) = (1/\sqrt{2\pi\sigma^2}) \exp(-(x_i-\mu)^2/(2\sigma^2))$.
Fatal Flaw: Purely text/URL-driven. Blind to image matrices (0% Quishing recall). Inability to query live DNS.

2.2 System B: Random Forest & Information Gain (Toolan & Carthy, 2010)

Constructs B bagged decision trees evaluating $m = \sqrt{d}$ feature splits minimizing Gini Impurity: $I_G(t) = 1 - \sum p_k^2$. Predictions aggregated via majority voting.
Fatal Flaw: High False Positive Rate on urgent business emails (FPR = 8.94%) because urgency keywords trigger phishing splits. Text trees cannot decode visual QR matrices (31.3% recall).

2.3 System C: Support Vector Machine (Basnet et al., 2008)

Maximizes margin hyperplane via dual Lagrangian: $\max_{\alpha} \sum \alpha_i - 0.5 \sum \alpha_i \alpha_j y_i y_j \exp(-\gamma \|x_i - x_j\|^2)$.
Fatal Flaw: Black-box scoring with $O(N^3)$ training cost, and cannot execute live asynchronous cryptographic queries.

3. Detailed Working Algorithm of SecureMail (Proposed)


SecureMail deploys a Tri-Modal Architecture uniting three orthogonal planes:

- **Modality 1 (Text NLP P_{nlp}):** Vectorizes 30+ engineered structural/intent features with linear-sigmoid calibration.
- **Modality 2 (Computer Vision Quishing P_{quish}):** Extracted DOM images -> Grayscale $I_{gray} = 0.299R + 0.587G + 0.114B$ -> Adaptive Gaussian thresholding -> pyzbar 2D finder pattern matrix decoding -> defanged URL analysis.
- **Modality 3 (Cryptographic DoH & D_{auth}):** Parallel DNS-over-HTTPS queries (<115ms) for authoritative SPF/DMARC alignment. Awards **Authentic Sender Discount ($D_{auth} = 0.22$)** when cryptographic proof is validated.
- **Bayesian Multimodal Fusion:** $P_{fused} = 1 - (1 - P_{nlp})(1 - P_{quish})$; $P_{final} = \text{clip}(P_{fused} + P_{penalty} - D_{auth}, 0, 1)$.


4. Quantitative Performance Benchmark (Test Set: N=1,500)

Model / Algorithm	Acc (%)	Prec (%)	Rec (%)	F1 (%)	FPR (%)	Quish Rec	Latency
Naive Bayes (Fette et al., 2007)	93.73	98.55	88.98	93.52	1.36	51.3%	0.01 ms
SVM (RBF Kernel) (Basnet et al., 2008)	92.47	99.54	85.56	92.03	0.41	49.3%	0.00 ms
Random Forest (Toolan & Carthy, 2010)	86.47	90.45	82.02	86.03	8.94	31.3%	0.00 ms
SecureMail Tri-Modal Hybrid (Proposed)	97.87	99.06	96.72	97.88	0.95	96.7%	4.12 ms


Empirical Benchmark Visualizations:




Ablation Study: Quantifying the Impact of Each SecureMail Mod



Training Progression & Feature Optimization Visualizations:



Calibrated Confusion Matrix (N=1,500 Test Partition):



5. Review Committee Viva Defense Cheat Sheet

Q: Why is your system taking 4.12 ms while Random Forest takes 0.00 ms?

Answer: In human perception, anything under 100 ms is instantaneous. Gmail takes 500--1,000 ms to render. Random Forest is 0 ms because it only does simple CPU math on text—it is blind to image quishing and causes false alarms. Spending 4.12 ms for computer vision detonation and cryptographic proof is an optimal security trade-off that raises accuracy from 86.47% to 97.87%.

Q: Why is Quishing Recall 96.7% and not 100%?

Answer: Our computer vision pipeline achieves 96.7% recall by extracting embedded QR matrices, applying adaptive contrast thresholding, and decoding finder patterns with pyzbar. The realistic ~3.3% missed margin reflects real-world adversarial edge cases (optical distortion, severe blur, or damaged finder patterns beyond Reed-Solomon parity recovery). Prior text-only baselines achieved only 31.3%--51.3%.

Q: How did you achieve a sub-1% False Positive Rate (0.95%) on urgent emails?

Answer: Via the **Authentic Sender Discount (D_auth = -22%)**. When real university or banking circulars use urgent words, our live DNS-over-HTTPS check validates their authentic SPF/DKIM/DMARC records and discounts the score, reducing false alarms by ~89% compared to Random Forest (8.94%).