

SecureMail: A Tri-Modal Hybrid Email Security Gateway and Forensics Framework for Zero-Trust Phishing and Optical Quishing Mitigation

Samir Khurshid

Enrollment: 2303031260196
2303031260196@paruluniversity.ac.in

M Harshith Chowdary

Enrollment: 2303031260149
2303031260149@paruluniversity.ac.in

Aman Kumar

Enrollment: 2303031260009
2303031260009@paruluniversity.ac.in

M. Prem Sai

Enrollment: 2303031260145
2303031260145@paruluniversity.ac.in

Department of Computer Science and Engineering (Cyber Security)
Parul Institute of Engineering & Technology, Parul University, Vadodara, Gujarat, India
Under the Guidance of: Prof. Akash Patel & Dr. Snehlata Barde

Abstract - Most enterprise email defenses still operate under a broken assumption: that malicious messages arrive as plain text. Modern attackers have moved on. They now routinely weaponize optical QR codes (quishing) and generate bespoke pretexts using large language models, slipping past perimeter secure email gateways without triggering a single lexical rule. When gateways do flag messages, they often misread urgent wording from legitimate internal senders, burying security teams under false alerts. To tackle this, we built SecureMail-a tri-modal defense framework that shifts inspection directly to the client reading pane. SecureMail couples three synchronized detection layers: (1) an event-driven Manifest V3 browser extension that captures DOM mutations in real time without storing user emails, (2) an optical computer vision module applying bilateral filtering and Otsu binarization to detonate embedded QR matrices, and (3) an asynchronous DNS-over-HTTPS resolver validating SPF, DKIM, DMARC, and ARC alignment over TLS 1.3. We also introduce the Authentic Sender Discount ($D_{auth} = 0.22$), a Bayesian probability modifier that cancels out urgency-driven false alarms when cryptographic sender proof is intact. Evaluated on an unseen 1,500-sample test partition drawn from an 8-corpus benchmark of 5,000 emails (including Nazario, Enron, and contemporary quishing campaigns), SecureMail Tier 1 (Quick Scan) achieves 97.87% accuracy, 98.4% quishing recall, and cuts the false positive rate on urgent institutional emails down to 0.95%-all within 110 ms of local execution time. For ambiguous edge cases, Tier 2 (Deep Scan) routes payload hashes to a multi-engine sandbox, reaching 99.8% detection across 87 commercial engines. We evaluate SecureMail against 12 baseline models spanning 2007 to 2026, establishing an effective operational blueprint for client-edge email defense.

Index Terms - Email Security, Optical Quishing, Computer Vision, DNS-over-HTTPS (DoH), Homoglyph Detection, Zero-Trust Architecture, Machine Learning, Transformer Language Models, Business Email Compromise (BEC), Forensic Ingestion.

I. INTRODUCTION

Enterprise security teams spend millions on perimeter email defenses, yet spear-phishing remains the top vector for ransomware and corporate data breaches [1], [2]. The reason is architectural. Traditional boundary Mail Transfer Agents (MTAs) and Secure Email Gateways (SEGs) inspect messages as they cross the network perimeter, relying on known IP blocklists, regex signatures, and statistical language heuristics [2], [3]. If a threat slips through or originates from a compromised internal mailbox, the perimeter gateway never sees how the message renders inside the employee's browser or desktop mail client [1].

Compounding this visibility gap, phishing campaigns between 2023 and 2026 adopted multimodal evasions that render text-only filters useless. The most aggressive is optical quishing-embedding phishing lures inside Quick Response (QR) codes [6], [17], [18]. In a recent empirical study across 71,000 corporate phishing messages, Zhang et al. [18] found that QR phishing achieved victim click rates matching standard link attacks, while sailing past commercial SEGs in 91.8% of deliveries. The explanation is simple: an optical QR matrix contains no readable ASCII strings in the raw MIME stream. Text models and keyword scrapers find nothing to analyze [17].

At the same time, attackers leverage generative AI to craft grammatically clean spear-phishing pretexts at scale, stripping away traditional grammatical flaws [21], [22]. They also register Internationalized Domain Names (IDNs) containing Cyrillic or Greek lookalikes to fool human eyes while bypassing strict ASCII string checks [23].

We built SecureMail to bridge these defensive blind spots. By pairing a lightweight Manifest V3 client extension with a local computer vision pipeline, asynchronous DNS-over-HTTPS (DoH) provenance checks, and a Bayesian trust discount, SecureMail inspects emails at the moment of rendering-stopping multimodal threats in under 110 ms without compromising inbox privacy.

II. PROBLEM STATEMENT AND RESEARCH OBJECTIVES

Our examination of contemporary email security revealed three recurring points of failure: 1) Client-Edge Invisibility: Boundary gateways evaluate messages before delivery, maintaining zero visibility over dynamic DOM rendering, delayed script payloads, or compromised internal tenant mailboxes [1], [24]. 2) Computer Vision Blindness: Mainstream filters parse text tokens, headers, and basic HTML markup. When a threat arrives as an image, text parsers see only raw base64 noise, missing 91.8% of QR phishing payloads [6], [17], [18]. 3) Alert Fatigue from Urgent Clean Mail: Spam models penalize urgent words ('immediate response required', 'payroll deadline'). Lacking cryptographic proof of sender identity, they misflag legitimate university, banking, and HR notices, causing false alarm rates up to 11.20% and training users to ignore security banners [8], [16].




Fig. 1. Threat Vector Surge Dynamics (2020-2026) and Problem-to-Objective Engineering Alignment Matrix for Zero-Trust Email Security.

Fig. 1 maps this empirical reality. In Fig. 1(a), threat volumes tracked from 2020 to 2026 demonstrate how traditional link phishing has plateaued, whereas generative pretexts (+2,100%) and optical quishing (+1,750%) have exploded. In Fig. 1(b), we contrast these evasion vectors against our engineering response matrix, highlighting how SecureMail eliminates the 91.8% quishing bypass and reins in the 11.2% false alarm rate.

To resolve these challenges, we defined six practical engineering goals:

- Objective 1 (Client-Edge DOM Ingestion): Deploy a Manifest V3 browser extension that hooks into Gmail and Outlook Web via MutationObserver, parsing inbound DOM structures locally without storing credentials or messages on disk.
- Objective 2 (Asynchronous Micro-Gateway): Build a Python FastAPI gateway delivering sub-110 ms local inference so users never experience reading-pane latency.
- Objective 3 (Cryptographic DoH Resolver): Implement an asynchronous DNS-over-HTTPS resolver querying Cloudflare (1.1.1.1) and Google (8.8.8.8) [7], [10] to verify SPF (RFC 7208) [12], DKIM (RFC 6376) [13], and DMARC (RFC 7489) [14] alignment over TLS 1.3 [11], alongside Authenticated Received Chain (ARC) provenance [15].
- Objective 4 (Optical Computer Vision Detonation): Construct a computer vision pipeline using bilateral filtering, Otsu binarization, and morphological dilation to decode degraded QR images in under 30 ms.
- Objective 5 (Authentic Sender Discount Formulation): Establish a Bayesian risk fusion formula incorporating our Authentic Sender Discount ($D_{auth} = 0.22$) to drive urgency false alarms on verified institutional emails below 1.0%.
- Objective 6 (Reproducible Multi-Model Benchmark): Benchmark SecureMail against 12 historical and contemporary models spanning 2007 through 2026 across standardized corpora (Nazario, Enron, PhishTank, SpamAssassin) with fully published artifacts.

III. THEORETICAL FOUNDATIONS AND THREAT MODELING

To establish mathematical rigor, SecureMail models email risk across four computational domains: Information Theory, Metric Space Distances, Image Matrix Transformations, and Bayesian Probability Fusion.

A. Information-Theoretic Entropy of Obfuscated URLs

Attackers frequently inject pseudo-random tracking strings and Domain Generation Algorithm (DGA) labels to dodge keyword matching. We gauge string randomness through Shannon Entropy $H(S)$. For any string S of length L with unique characters c_1 through c_k and empirical probability $p(c_i) = \text{count}(c_i) / L$:

$$H(S) = - \sum_{i=1}^k p(c_i) \log_2 p(c_i) \quad (1)$$

Established enterprise domains usually exhibit entropy between 2.4 and 3.6. In contrast, randomized tracking tokens and phishing hashes regularly exceed $H(S) > 4.25$. SecureMail evaluates $H(S)$ across domain labels and path components, assigning a proportional risk penalty $P_{entropy}$ between 0 and 0.25.

B. Metric Space Distance for IDN Homoglyph Detection

Phishing operators frequently register Internationalized Domain Names (IDNs) using Cyrillic or Greek lookalikes (such as Cyrillic 'a' [U+0430] matching Latin 'a' [U+0061]) to impersonate brands. We quantify visual similarity using Normalized Levenshtein Distance. Let $\text{Lev}(D_1, D_2)$ denote the minimum edit operations converting D_1 into brand target D_2 :

$$\text{Sim}_{\text{Lev}}(D_1, D_2) = 1 - \text{Lev}(D_1, D_2) / \max(|D_1|, |D_2|) \quad (2)$$

Before computing string distance, our engine applies Unicode Confusable Normalization (UTR #39) [28], mapping international glyphs to canonical Latin skeletons. If $\text{Sim}_{\text{Lev}}(D_{\text{confusable}}, D_{\text{brand}}) \geq 0.85$ while $D_1 \neq D_2$, SecureMail flags an IDN masquerade alert with penalty $P_{homoglyph} = 0.80$.

C. Computer Vision Image Matrix Filtering & Otsu Criterion

Embedded QR codes in emails are frequently degraded by aggressive compression, intentional contrast lowering, or background patterns. We model each extracted image as a 2D intensity grid $I(x, y)$. Instead of standard Gaussian blur, which rounds off finder corners, we use Bilateral Filtering to smooth noise while preserving critical high-frequency edges:

$$I_{\text{filtered}}(p) = (1 / W_p) \sum_{q \text{ in } \Omega(p)} I(q) f_r(\|I(q) - I(p)\|) g_s(\|q - p\|) \quad (3)$$

where W_p normalizes weights, g_s is a spatial Gaussian kernel, and f_r is an intensity radiometric kernel. To binarize the image without manual threshold tuning, our engine maximizes Otsu's inter-class variance $\sigma_B^2(t)$:

$$\sigma_B^2(t^*) = \max_t [\omega_0(t) \omega_1(t) (\mu_0(t) - \mu_1(t))^2] \quad (4)$$

where ω_0 and ω_1 represent background and foreground class probabilities at threshold t , and μ_0 , μ_1 represent their respective mean intensities. Morphological dilation with a 3x3 structuring element then joins broken timing patterns, enabling pyzbar to decode payload URLs with 98.4% recall.

D. Cryptographic DNS Provenance & Authentic Sender Discount

To resolve the chronic friction between urgent email text and legitimate administrative notices, we integrate Forward-Confirmed Reverse DNS (FCrDNS) and DNS-over-HTTPS (DoH) header verification into our Bayesian fusion:

$$P_{\text{fused}} = 1 - (1 - P_{\text{nlp}}) * (1 - P_{\text{quish}}) \quad (5)$$

$$P_{\text{final}} = \min(1.0, \max(0.0, P_{\text{fused}} + P_{\text{penalty}} - D_{\text{auth}} * I_{\text{auth}})) \quad (6)$$

Here, I_{auth} is a strict binary indicator returning 1 only when SPF, DKIM, and DMARC all return verified 'pass' states over secure DoH. Through tuning across 1,500 validation emails, we established an

optimal discount factor $D_{auth} = 0.22$, effectively neutralizing urgency alarms on legitimate, cryptographically authentic traffic.

IV. RELATED WORK & EVOLUTION OF PHISHING DETECTION (2007–2026)

Over the past twenty years, email security research evolved across three major phases: handcrafted feature engineering (2007-2015), deep semantic modeling (2016-2022), and multimodal zero-trust architectures (2023-2026).

A. Classical Classifiers (2007-2015): In 2007, Fette et al. [2] introduced PILFER, pioneering machine learning on the Nazario corpus [5] using ten handcrafted lexical features across Naive Bayes and Random Forest. Toolan and Carthy [3] expanded the feature vector to 21 structural attributes, finding Random Forest provided greater stability under skewed token distributions. Basnet, Mukkamala, and Sung [4] trained Support Vector Machines (SVM) with RBF kernels, showing high accuracy on pristine text but sharp degradation on noisy vocabularies. However, these early systems evaluated ASCII tokens in isolation; they could neither decode visual QR payloads nor verify cryptographic DNS headers [16].

B. The Multimodal Evasion Era (2023-2024): Vasek, Wadleigh, and Christin [6] conducted the first comprehensive wild empirical study of optical quishing, demonstrating that QR deception slipped past commercial perimeter gateways in over 85% of campaigns. Foster, Johnson, and Lee [8] audited SPF, DKIM, and DMARC adoption across 100,000 domains, uncovering widespread misconfigurations that permitted trivial sender impersonation. Chen, Zhang, and Wang [9] proposed visual phishing detection with deep embeddings, though their pipeline required computationally expensive GPU infrastructure. Concurrently, Zhang, Shen, and Ren [23] characterized the exploitation of Internationalized Domain Names (IDNs), proving that visual lookalikes reliably defeated ASCII exact-match filters across modern registries.

C. Recent Breakthroughs (2025-2026): In 2025, Zhang, Wang, and Christin [18] analyzed 71,000 QR phishing campaigns, confirming that quishing achieved click rates comparable to classical link phishing while evading 91.8% of SEGs. Martinez and Alvarez [19] investigated dual-modulated QR codes for physical authentication. Sharma, Singh, and Rastogi [20] combined DistilBERT, CNN-LSTM, and explainable AI, achieving 98.2% accuracy on text corpora. Nguyen, Tran, and Dubois [22] built stylometric transformers to detect generative AI phishing, while Al-Hubaishi, Mohammad, and Saeed [25] introduced a dual-path deep neural net for BEC header provenance. In 2025, Carle, Valenza, and Trevisan [26] confirmed that combining forward-confirmed reverse DNS with DoH was crucial for reliable automated gateway verification. Most recently in 2026, Park, Lee, and Kim [17] introduced QuishFormer, a vision transformer tailored for QR extraction; Roy, Das, and Chatterjee [21] presented E-PhishGEN to benchmark synthetic LLM campaigns; and Patel, Joshi, and Shah [24] demonstrated zero-trust threat intelligence integration at the gateway layer.

TABLE I: Taxonomy of State-of-the-Art Phishing Detection Literature (2007-2026)

Study / Reference	Year	Core Architecture / Algorithm	Input Modality	Quishing?	DoH DNS?
Fette et al. [2]	2007	Naive Bayes & Decision Trees (PILFER)	Raw ASCII Text Tokens	No (0%)	No
Basnet et al. [4]	2008	Support Vector Machines (RBF Kernel)	Handcrafted Lexical Vectors	No (0%)	No
Toolan & Carthy [3]	2010	Random Forest (21 Selected Features)	Structural Email Metadata	No (0%)	No
Sharma et al. [20]	2025	Hybrid DistilBERT + CNN-LSTM + XAI	Contextual Text Embeddings	No (12%)	No
Nguyen et al. [22]	2025	Stylometric	Deep Syntactic	No (14%)	No

		RoBERTa Transformer	Embeddings		
Park et al. [17]	2026	QuishFormer Vision Transformer (ViT)	Visual Image Matrices	Yes (97.2%)	No
SecureMail (Ours)	2026	Tri-Modal Hybrid (CV + DoH + Bayesian)	Text + QR Vision + DNS Provenance	Yes (98.4%)	Yes (DoH)

Table I highlights the central architectural divide in literature: prior works either focused purely on text tokens while ignoring visual quishing, or proposed heavy vision networks that omitted cryptographic header validation. SecureMail is the first architecture to harmonize all three modalities into a single, low-latency client-edge system.

V. PROPOSED SECUREMAIL SYSTEM ARCHITECTURE

SecureMail organizes its defense into four loosely coupled tiers: (1) Client-Edge Layer, operating an event-driven Manifest V3 extension with MutationObserver hooks inside Gmail and Outlook Web; (2) Security Gateway Layer, written in Python 3.10+ using FastAPI and an asynchronous Uvicorn ASGI server; (3) Analysis Engine Layer, coordinating concurrent workers for DoH queries, OpenCV QR decoding, Levenshtein distance, and Threat Vault caching; and (4) Persistence & Audit Layer, logging cryptographically signed SQLite records, enforcing volatile memory boundaries, and generating ISO/IEC 27001 aligned forensic PDF reports.




Fig. 2. Secure Extension-to-Gateway Trust Separation and Volatile Memory Boundary.

Fig. 2 illustrates our extension-to-gateway trust separation and memory lifecycle. By isolating DOM parsing within the browser sandbox and streaming ephemeral JSON payloads over mutual TLS to FastAPI, SecureMail prevents session token leakage while ensuring raw message bodies are purged from RAM immediately after feature scoring.

TABLE II: Documented Technology Stack and Component Allocation

Subsystem / Layer	Technology Component	Engineering Justification & Performance Standard
Client-Edge Extension	Chrome Manifest V3 / WebExtensions	Zero persistent email storage, non-intrusive DOM extraction (<15 ms)
REST API Gateway	FastAPI / Uvicorn ASGI / Python 3.10+	Asynchronous non-blocking concurrency, Pydantic data validation
Machine Learning Engine	Scikit-Learn / TF-IDF Vectorizer	Local lightweight inference (<15 ms), zero third-party token leakage
Computer Vision Pipeline	OpenCV 4.8 / NumPy / pyzbar	Adaptive thresholding and morphological dilation for QR detonation
Cryptographic DNS Layer	https / Cloudflare & Google DoH	Asynchronous RFC 8484 DNS-over-HTTPS resolution over TLS 1.3
Threat Intelligence Vault	VirusTotal v3 REST API /	Multi-engine sandbox

	SQLite	detonation (87+ scanners) for Tier 2 Deep Scan
Audit & Forensics Engine	ReportLab / PyCryptodome (AES-256)	Cryptographically signed tamper-evident digital forensic reports

Table II outlines the engineering rationale behind our component choices, balancing strict zero-trust privacy boundaries with a median client-edge execution latency under 110 ms.

VI. MULTIMODAL DETECTION METHODOLOGY & FEATURE ENGINEERING

Our detection pipeline processes inbound emails across three concurrent, synchronized inspection tracks: 1) Stylometric & NLP Track: The engine sanitizes text extracted from DOM nodes, extracting unigram and bigram TF-IDF weights, URL Shannon entropy, punctuation density, and urgency token frequencies to produce preliminary risk estimate P_{nlp} .

Optical Quishing (QR Code) Inspection Flow




Fig. 3. Optical Quishing (QR Code) Inspection and Payload Decoding Flow.

Fig. 3 illustrates our 4-stage computer vision detonation workflow. When incoming messages contain embedded images or inline attachments, OpenCV applies bilateral filtering and Otsu thresholding to reconstruct degraded QR matrices, allowing pyzbar to decode payload URIs with 98.4% recall.

SecureMail End-to-End Threat Detection Flow




Fig. 4. End-to-End Threat Detection and Forensic Ingestion Flow.

Fig. 4 tracks our end-to-end asynchronous processing pipeline, illustrating how telemetry moves from client DOM capture through parallel worker evaluation, Bayesian fusion, and automated PDF forensic reporting.

Real-Time DoH Cryptographic DNS Verification




Fig. 5. Real-Time DNS-over-HTTPS (DoH) Authentication Verification Flow.

Fig. 5 details the real-time DoH verification sequence. Asynchronous TLS 1.3 queries [11] to Cloudflare and Google DoH endpoints [7] validate SPF [12], DKIM [13], and DMARC [14] alignment, as well as intermediate ARC hops [15], triggering the Authentic Sender Discount ($D_{auth} = 0.22$) to suppress urgency false alarms.

SecureMail Security Testing Pipeline




Fig. 6. Multi-Layer Security Hardening and Vulnerability Assessment Pipeline.

Fig. 6 details the multi-layer security hardening and vulnerability assessment pipeline, enforcing strict Content Security Policies (CSP), CORS controls, and volatile memory scrubbing to eliminate attack surfaces on the gateway itself.

Defect Reproduction & Regression Lifecycle




Fig. 7. Regression Testing Lifecycle and Webmail DOM Evolution Harness.

Fig. 7 outlines the regression testing lifecycle and webmail DOM evolution harness, ensuring resilient selector parsing against dynamic DOM mutations introduced by webmail service updates.

VII. EXPERIMENTAL SETUP, BENCHMARK DATASETS & HARDWARE HARNESS

All evaluations were conducted on an AMD Ryzen 7 5800H system (8 cores, 16 threads @ 3.2 GHz) with 32 GB DDR4 RAM, running across Windows 11 and Ubuntu 22.04 LTS environments to ensure reproducible execution.

TABLE III: Benchmark Datasets Characteristics and Modality Distribution

Dataset Name	Source / Provider	Samples	Benign	Phishing	Primary Threat Modality
Nazario Corpus [5]	Jose Nazario	4,500	0	4,500	Classical Link & Credential Phishing
Enron Corpus [27]	CMU / FERC	5,000	5,000	0	Clean Corporate Institutional Mail
PhishTank Feed	OpenDNS / Cisco	3,200	0	3,200	Zero-Hour Verified Malicious URLs
SpamAssassin	Apache Foundation	6,000	4,150	1,850	Heuristic Spam & Lexical Phishing
IWSPA Task	ACM IWSPA	2,800	1,400	1,400	Targeted Spear-Phishing Pretexting
Custom Quishing	Empirical (2024-26)	1,250	250	1,000	Visual Optical QR Code Payloads
IDN Homoglyph	Unicode TR	850	350	500	Cyrillic/Greek Lookalike

[28]	#39	Domains			
Generative AI	Synthesized (LLM)	1,000	0	1,000	LLM Contextual Spear-Phishing

Table III summarizes the characteristics and sample distributions of our 8 benchmark datasets, incorporating legitimate corporate communications from the Enron corpus [27] alongside internationalized spoofing sets categorized under Unicode TR #39 [28]. From this master corpus, a stratified evaluation partition of 5,000 samples was drawn: 3,500 samples (70%) for training, 500 samples (10%) for validation, and 1,500 samples (20%) held out as a strictly unseen benchmark test partition.

TABLE IV: Mathematical Feature Formulations and Information-Theoretic Parameters

Feature Identifier	Mathematical Formulation	Detection Purpose & Parameter Bounds
Shannon Entropy $H(S)$	$H(S) = -\sum p(c_i) \log_2 p(c_i)$	Detects randomized URL strings and DGA domains (Threshold: >4.25)
Levenshtein Sim	$Sim_Lev = 1 - Lev / max_len$	Identifies brand homoglyphs and punycode masquerading (Threshold: ≥ 0.85)
Otsu Binarization	$\sigma_B^2 = \max [w_0 w_1 (u_0 - u_1)^2]$	Maximizes contrast between QR finder patterns and background noise
Bilateral Filter	$I_{filt} = (1/W_p) \sum I_{f_r} g_s$	Smooths gradient artifacts while preserving critical module borders
Authentic Discount	$D_{auth} = 0.22 * I(DMARC \wedge DKIM)$	Neutralizes urgency score on cryptographically proven sender domains
Fused Probability	$P_{fused} = 1 - (1 - P_{nlp})(1 - P_{quish})$	Combines orthogonal textual and visual risk indicators into joint posterior

The mathematical feature formulations, bounds, and detection purposes are summarized in Table IV, establishing the quantitative inputs for our Bayesian inference engine.

VIII. SYSTEM VERIFICATION & AUTOMATED QUALITY ASSURANCE EVIDENCE

Before running benchmark evaluations, we validated the entire SecureMail codebase through 123 automated test suites across eight operational domains.




Fig. 8. SecureMail System Testing Architecture and Multi-Layer Operational Verification Boundaries.

Fig. 8 visualizes the system testing architecture and multi-layer operational verification boundaries, orchestrating our 123 automated test suites across unit, integration, computer vision, DoH resolver, and security layers.

TABLE V: Testing Categories and Automated Evidence Status (123 Automated Tests)

Testing Category	Verification Scope	Executed Tests	QA Status
Unit Testing	Tokenization, regex parsing, TF-IDF vectorization	32 Tests	100% Passed
Integration Testing	FastAPI routing, SQLite async pool, session handling	24 Tests	100% Passed
Computer Vision Testing	Otsu thresholding, bilateral smoothing, QR extraction	18 Tests	100% Passed
DoH Cryptographic Testing	Cloudflare/Google DoH DNSSEC, SPF/DKIM verification	14 Tests	100% Passed
Security Hardening Testing	CORS policy, JWT expiration, volatile memory purge	12 Tests	100% Passed
Performance & Stress	Concurrent locust loads (500 req/sec), memory leak profile	9 Tests	100% Passed
Browser Extension E2E	Manifest V3 DOM MutationObserver in Gmail/Outlook	8 Tests	100% Passed
Forensics Audit Validation	ReportLab PDF generation, AES-256 integrity hash	6 Tests	100% Passed

Table V summarizes the verification scope and execution status across all 123 automated test suites, confirming 100% pass rates prior to benchmark ingestion.

IX. COMPARATIVE ALGORITHMIC BENCHMARKS & MULTI-MODEL EVALUATION

We benchmarked SecureMail against 12 baseline models spanning traditional machine learning, gradient boosting, and deep transformer architectures. As depicted in the Receiver Operating Characteristic (ROC) curves in Fig. 9, SecureMail achieves an Area Under the Curve (AUC) of 0.999, outperforming Naive Bayes (AUC = 0.961), SVM (AUC = 0.954), and Random Forest (AUC = 0.912).




Fig. 9. Receiver Operating Characteristic (ROC) Benchmark Curves (AUC = 0.999).




Fig. 10. Optical Quishing Recall and False Positive Rate across Benchmark Models.




Fig. 11. False Positive Rate on High-Urgency Benign Institutional Emails (N=738) Contrasting Classical Classifiers against SecureMail Provenance Discount.

Fig. 10 and Fig. 11 highlight the severe performance collapse exhibited by traditional models when confronted with multimodal attacks. In optical quishing (Fig. 10), classical models achieve exactly 0.0% recall because the attack payload resides in pixel space. On high-urgency benign emails (Fig. 11, N=738), Naive Bayes and Random Forest generate 11.20% and 8.94% false positive rates respectively, whereas SecureMail achieves 0.95% FPR through its authentic provenance discount.

TABLE VI: Baseline Algorithmic Comparison across Traditional Classifiers (N=1,500 Test Samples)

Model / Baseline Algorithm	Training Corpus	Accuracy (%)	Precision (%)	Recall (%)	F1-Score (%)	Urgency FPR (%)
Naive Bayes (PILFER) [2]	Nazario + Enron	93.52	91.80	94.20	92.98	11.20
SVM (RBF Kernel) [4]	Nazario + Enron	92.03	90.15	93.10	91.60	12.80
Random Forest (Toolan) [3]	Nazario + Enron	86.03	84.10	87.20	85.62	8.94
Decision Tree (C4.5 Baseline)	Nazario + Enron	83.15	81.20	84.50	82.82	14.50
Logistic Regression (L2)	Nazario + Enron	88.40	87.10	89.20	88.14	10.30

Table VI benchmarks traditional machine learning baselines across the 1,500 unseen test partition, demonstrating that classical approaches struggle with vocabulary drift and urgency false alarms.

TABLE VII: Advanced Gradient Boosting & Transformer Benchmark Comparison

Advanced Model Architecture	Publication Year	Accuracy (%)	Precision (%)	Recall (%)	F1-Score (%)	Latency (ms)
XGBoost Classifier	2021	94.85	93.90	95.10	94.50	42 ms
CatBoost Classifier	2022	95.20	94.40	95.60	95.00	48 ms
Hybrid DistilBERT + CNN-LSTM [20]	2025	96.20	95.40	96.80	96.09	160 ms
DistilBERT Transformer [20]	2025	96.50	95.80	96.90	96.35	145 ms
Stylometric RoBERTa [22]	2025	96.85	96.20	97.10	96.65	182 ms
QuishFormer ViT [17]	2026	94.10	93.50	94.60	94.05	240 ms

Table VII compares advanced gradient boosting classifiers and deep transformer language models, illustrating higher textual accuracy at the expense of significant computational inference overhead.

X. MASTER EMPIRICAL RESULTS: RESULT TABLE 9 & IN-DEPTH DISCUSSION

To synthesize performance across two decades of email defense, Table IX compares 12 distinct models across eight standardized metrics on our held-out test partition (N=1,500).

TABLE VIII: Modality Ablation Performance Degradation Matrix

Ablated System Configuration	Excluded Component	Accuracy (%)	F1-Score (%)	Quishing Rec (%)	Urgency FPR (%)
Full SecureMail Tri-Modal (Tier 1)	None (Complete System)	97.87	97.88	98.4	0.95
Ablation A: Without Computer Vision	OpenCV & pyzbar Detonation	91.24 (-6.63)	90.85 (-7.03)	0.0 (-98.4)	1.12 (+0.17)
Ablation B: Without	DoH Resolver &	93.40 (-)	93.12 (-)	98.4 (0.0)	8.75


TABLE IX: Master Multi-Model Benchmark Evaluation and Performance Comparison (2007-2026 Baselines vs. SecureMail)

Model / Architecture	Year	Input Modality	Acc (%)	Prec (%)	Rec (%)	F1 (%)	Urg FPR (%)	Quish Rec (%)	Latency (ms)
Naive Bayes (PILFER) [2]	2007	ASCII Text Tokens	93.52	91.80	94.20	92.98	11.20	0.0	18 ms
SVM (RBF Kernel) [4]	2008	Lexical Text Vectors	92.03	90.15	93.10	91.60	12.80	0.0	24 ms
Random Forest [3]	2010	21 Lexical Features	86.03	84.10	87.20	85.62	8.94	0.0	31 ms
Decision Tree (C4.5)	2012	Syntactic Token Splits	83.15	81.20	84.50	82.82	14.50	0.0	12 ms
Logistic Regression	2014	L2 Token Probabilities	88.40	87.10	89.20	88.14	10.30	0.0	15 ms
XGBoost Classifier	2021	Gradient Boosted Trees	94.85	93.90	95.10	94.50	6.80	0.0	42 ms
CatBoost Classifier	2022	Categorical Boosting	95.20	94.40	95.60	95.00	6.20	0.0	48 ms
DistilBERT [20]	2025	Contextual Embeddings	96.50	95.80	96.90	96.35	4.80	12.0	145 ms
RoBERTa Stylometric [22]	2025	Deep Syntactic Attention	96.85	96.20	97.10	96.65	4.20	14.5	182 ms
Hybrid CNN-LSTM [20]	2025	Spatial-Temporal Tokens	96.20	95.40	96.80	96.09	5.10	8.0	160 ms
QuishFormer ViT [17]	2026	Visual Patch Matrices	94.10	93.50	94.60	94.05	7.90	97.2	240 ms
Dual-Path BEC [25]	2025	Header Deep Neural Net	95.60	95.10	96.00	95.55	3.50	0.0	95 ms
SecureMail Tier 1 (Quick)	2026	Tri-Modal Hybrid	97.87	97.82	97.94	97.88	0.95	98.4	110 ms
SecureMail Tier 2 (Deep)	2026	Hybrid + Cloud Sandbox	99.80	99.75	99.85	99.80	0.15	99.6	4,200 ms

Analysis of Table IX demonstrates that SecureMail Tier 1 establishes state-of-the-art performance among real-time local gateways. While standalone transformers (DistilBERT [20], RoBERTa [22]) achieve competitive lexical accuracy (96.50%-96.85%), their inability to evaluate visual matrices results in catastrophic failure against quishing (12.0%-14.5% recall). Conversely, vision transformers like QuishFormer [17] achieve 97.2% quishing recall but lack cryptographic DNS validation, suffering an elevated urgency false alarm rate of 7.90% and higher inference latency (240 ms). SecureMail Tier 1 harmonizes all modalities into a coherent probability distribution, achieving 97.87% accuracy and 98.4% quishing recall with an exceptional 0.95% urgency FPR at 110 ms latency.

XI. ABLATION STUDIES, PRECISION-RECALL DYNAMICS & OPERATIONAL TRADE-OFFS

To quantify the individual contribution of each subsystem, we performed granular ablation profiling across our entire evaluation partition.



Cryptographic DoH	D_auth Discount	4.47)	4.76)	(+7.80)
Ablation C: Without Homoglyph Engine	Levenshtein & Unicode Tables	94.80 (-3.07)	94.62 (-3.26)	98.4 (0.0) 1.20 (+0.25)
Ablation D: Text-Only Minimal Model	Vision, DoH & Homoglyphs	89.15 (-8.72)	88.70 (-9.18)	0.0 (-98.4) 10.15 (+9.20)

Table VIII details the performance degradation observed under systematic modality ablation. Disabling the computer vision pipeline causes quishing recall to collapse to 0.0%, while disabling DoH provenance causes the urgency false positive rate to spike from 0.95% to 8.75%.

Fig. 12. Ablation Study Quantifying the Impact of Individual SecureMail Modalities.

Fig. 12 illustrates the ablation study quantifying the marginal contribution of each modality on overall detection accuracy and recall.




Fig. 13. Feature Correlation Heatmap Quantifying Attack Markers and Provenance Discounts.

Fig. 13 presents the feature correlation heatmap across attack indicators, confirming that homoglyph, quishing, and entropy features correlate

strongly with ground-truth phishing labels while remaining orthogonal to benign metadata.




Fig. 14. Multi-Metric Model Convergence Trajectory Across Sequential Architectural Modalities.

Fig. 14 plots the multi-metric model convergence trajectory across sequential architectural modalities, showing rapid convergence and stable loss stabilization within 15 epochs.




Fig. 15. Precision-Recall Curves Across Evasive Threat Modalities Illustrating High Positive Predictive Value at Extended Sensitivity Thresholds.

Fig. 15 benchmarks the Precision-Recall dynamics across optical quishing (AP = 0.994), weaponized attachments (AP = 0.992), IDN homoglyphs (AP = 0.988), and BEC (AP = 0.976), confirming high positive predictive value across diverse operating thresholds.




Fig. 16. Multi-Vector Threat Detection Coverage Benchmark Contrasting SecureMail against Contemporary and Classical Baselines across Six Attack Vectors.

Fig. 16 presents a comprehensive radar comparison across six primary attack vectors, illustrating SecureMail's dominant and balanced coverage exceeding 96.7% across all threat vectors.




Fig. 17. Confusion Matrix on Benchmark Test Partition (N=1,500).

Fig. 17 displays the calibrated confusion matrix across 1,500 unseen evaluation emails, demonstrating only 16 false positives and 16 false negatives.




Fig. 18. Subsystem Latency Profiling across Quick Scan Pipeline ($N=1,500$) Comparing Median (p50) and 95th Percentile (p95) Execution Latencies.




Fig. 19. Operational Pareto Frontier Quantifying Empirical Detection Depth versus End-to-End Processing Latency.

Fig. 18 illustrates the empirical latency distribution across each subsystem of the Quick Scan pipeline over 1,500 evaluations. The median execution time (p50) is dominated by DoH DNS resolution (43.2 ms) and OpenCV computer vision decoding (27.5 ms), followed by Scikit-Learn TF-IDF inference (14.6 ms), client DOM parsing (11.8 ms), IDN homoglyph matching (7.9 ms), and Bayesian scoring (3.8 ms). Even at the 95th percentile (p95), total local inspection completes in approximately 168.1 ms, satisfying real-time browser extension constraints (<200 ms) without requiring external threat feed network hops. Fig. 19 further illustrates how SecureMail establishes a new operational Pareto frontier, offering 97.87% accuracy at 110 ms locally and 99.80% in deep sandbox mode.

XII. SYSTEM IMPLEMENTATION, OPEN DATASETS, AND WEB VERIFICATION LINKS

To support open science and scientific validation, all software modules, model weights, synthetic attack datasets, and interactive dashboards are documented and accessible through verified endpoints:

- 1) Project Source Code & Extension: The complete open-source codebase, comprising the Manifest V3 browser extension, FastAPI gateway, and automated unit tests, is available on GitHub: <https://github.com/samirkhurshid/SecureMail>.
- 2) Interactive Web Gateway Interface: The production webmail simulation client and real-time administrative threat monitoring dashboard operate locally on host port 5173 (<http://localhost:5173>).
- 3) Gateway REST API Documentation: Automated interactive OpenAPI/Swagger specification endpoints are exposed directly at <http://localhost:8000/docs> for external SIEM and SOAR forensic integration.
- 4) Cryptographic DoH Resolvers: Real-time DNS verification routes RFC 8484 queries to Cloudflare Primary (<https://cloudflare-dns.com/dns-query>) and Google Public DNS (<https://dns.google/resolve>) over TLS 1.3 [10].
- 5) Threat Intelligence Engine: Tier 2 multi-engine sandbox inspection integrates with the VirusTotal v3 REST API (<https://developers.virustotal.com/reference/overview>) across 87+ commercial antivirus engines [24].
- 6) Public Benchmark Data Repositories: Standardized evaluation datasets utilized in this study are permanently hosted at: Nazario Phishing Corpus (<https://monkey.org/~jose/phishing/>), CMU Enron Archive (<https://www.cs.cmu.edu/~enron/>), PhishTank (<https://phishtank.org/>), and Unicode Security TR#39 Confusables (<https://www.unicode.org/Public/security/latest/confusables.txt>).

XIII. LIMITATIONS AND ADVERSARIAL RESILIENCE CONSIDERATIONS

While SecureMail improves upon traditional boundary defenses, several operational and threat model limitations must be noted: 1) Adversarial QR Distortion: Adversaries may apply intentional micro-perturbations or fragmented dot matrices that fall below Otsu thresholding tolerances, requiring future integration of lightweight neural super-resolution models [17], [19]; 2) Delayed Ephemeral Redirects: Malicious QR codes can point to intermediate domains that serve benign pages during automated scanner inspection, redirecting to credential harvesters only after user interaction or specific mobile user-agent headers [18]; 3) Compromised Trusted Accounts: In Business Email Compromise (BEC) scenarios where an attacker sends mail from an authentic corporate mailbox passing SPF, DKIM, and DMARC, the Authentic Sender Discount D_{auth} will apply, placing the defense burden on stylometric NLP safeguards [25]; and 4) Client Hardware Load: Rapid continuous inbox scrolling on low-spec client machines can cause brief CPU spikes (3-5%) from DOM mutation observers, which we counter through 250 ms event debouncing.

XIV. FUTURE WORK

Based on our empirical results, we outline five key engineering directions: 1) On-Device Quantized Transformers: Embedding quantized DistilBERT or RoBERTa models (INT8/ONNX) directly within the browser runtime to counter generative LLM phishing [20], [21]; 2) Native SMTP Milter Integration: Developing high-throughput C/Python Milter daemons for Postfix and Microsoft Exchange gateways; 3) Sandboxed Detonation: Integrating microVM containers for zero-day attachment behavioral analysis [24]; 4) Consortium Threat Sharing: Evaluating privacy-preserving threat intelligence sharing across academic institutions; and 5) Browser Portability: Expanding support from Chromium Manifest V3 to Firefox WebExtensions and Apple Safari.

XV. CONCLUSION

In this work, we developed and validated SecureMail—a tri-modal zero-trust email security gateway designed to neutralize modern optical quishing, generative AI lures, IDN homoglyphs, and perimeter visibility gaps. By fusing client-edge DOM mutation capture with automated OpenCV computer vision QR decoding and TLS 1.3 DNS-over-HTTPS provenance verification, SecureMail overcomes the fundamental blind spots of text-only filters and legacy SEGs. The Authentic Sender Discount ($D_{auth} = 0.22$) eliminates alert fatigue by driving false alarms on urgent benign messages down to 0.95%. Evaluated against a 5,000-sample benchmark corpus, SecureMail Tier 1 delivers 97.87% accuracy and 98.4% quishing recall with an average local execution latency of 110 ms, while Tier 2 Deep Scan achieves 99.8% detection across 87 antivirus engines. The system offers a practical, computationally lightweight, and mathematically grounded blueprint for next-generation enterprise email defense.

REFERENCES

- [1] P. Kumar, S. Gupta, and B. Panda, "A comprehensive survey on email spoofing, phishing attacks, and defense mechanisms," *IEEE Communications Surveys & Tutorials*, vol. 24, no. 2, pp. 1120-1151, 2022.
- [2] I. Fette, N. Sadeh, and A. Tomasic, "Learning to detect phishing emails," in *Proc. 16th Int. Conf. on World Wide Web (WWW)*, 2007, pp. 649-656.
- [3] F. Toolan and J. Carthy, "Feature selection for machine learning in email phishing detection," *Computers & Security*, vol. 29, no. 7, pp. 807-814, 2010.
- [4] R. Basnet, S. Mukkamala, and A. H. Sung, "Detection of phishing attacks: A machine learning approach," in *Proc. Int. Conf. on Artificial Intelligence (ICAI)*, 2008, pp. 373-379.
- [5] J. Nazario, "Phishing corpus: Public collection of phishing emails," <http://monkey.org/~jose/phishing/>, 2006.
- [6] M. Vasek, J. Wadleigh, and N. Christin, "Quishing in the wild: Characterizing QR code phishing campaigns and optical evasion techniques," in *Proc. ACM SIGSAC Conf. on Computer and Communications Security (CCS)*, 2023, pp. 889-903.
- [7] E. Osterweil, M. Larson, and D. Massey, "DNS-over-HTTPS (DoH) in practice: Performance, privacy, and security implications," *IEEE Internet Computing*, vol. 26, no. 1, pp. 48-57, 2022.

- [8] M. Foster, R. Johnson, and H. Lee, "Empirical evaluation of DMARC, SPF, and DKIM adoption across enterprise networks," *IEEE Transactions on Information Forensics and Security*, vol. 18, pp. 2410-2423, 2023.
- [9] H. Chen, T. Zhang, and L. Wang, "Optical phishing detection via computer vision and deep feature embedding," *IEEE Access*, vol. 12, pp. 45120-45134, 2024.
- [10] P. Hoffman and P. McManus, "DNS Queries over HTTPS (DoH)," *IETF RFC 8484*, 2018.
- [11] E. Rescorla, "The Transport Layer Security (TLS) Protocol Version 1.3," *IETF RFC 8446*, 2018.
- [12] S. Kitterman, "Sender Policy Framework (SPF) for Authorizing Use of Domains in Email," *IETF RFC 7208*, 2014.
- [13] D. Crocker, T. Hansen, and M. Kucherawy, "DomainKeys Identified Mail (DKIM) Signatures," *IETF RFC 6376*, 2011.
- [14] M. Kucherawy and E. Zwicky, "Domain-based Message Authentication, Reporting, and Conformance (DMARC)," *IETF RFC 7489*, 2015.
- [15] K. Kucherawy, J. Levine, and M. Kucherawy, "Authenticated Received Chain (ARC) Protocol," *IETF RFC 8617*, 2019.
- [16] M. A. Al-Daeef, N. Basir, and M. Saudi, "Machine learning approaches for phishing email detection: A comprehensive benchmark," *IEEE Access*, vol. 9, pp. 168541-168553, 2021.
- [17] J. Park, K. Lee, and S. Kim, "QuishFormer: Multi-format transformer framework for heterogeneous QR phishing detection," *Journal of Information Security and Applications (Elsevier)*, vol. 88, p. 103982, 2026.
- [18] L. Zhang, Y. Wang, and M. Christin, "Large-scale empirical analysis of QR code phishing campaigns and user susceptibility," in *Proc. 20th ACM ASIA Conf. on Computer and Communications Security (ASIA CCS)*, 2025, pp. 412-426.
- [19] R. V. Martinez and D. Alvarez, "Self-authenticating dual-modulated QR codes for phishing prevention," *IEEE Security & Privacy*, vol. 23, no. 1, pp. 34-43, 2025.
- [20] A. Sharma, P. K. Singh, and V. Rastogi, "Context-aware hybrid framework integrating DistilBERT, CNN-LSTM, and explainable AI for sophisticated phishing detection," in *Proc. 17th IEEE Int. Conf. on Contemporary Computing (IC3)*, 2025, pp. 112-119.
- [21] S. Roy, N. Das, and K. Chatterjee, "E-PhishGEN: Evaluating and mitigating generative large language model phishing attacks," in *Proc. ACM Conf. on Computer and Communications Security (CCS)*, 2026, pp. 581-596.
- [22] T. Nguyen, H. Tran, and E. Dubois, "Stylometric and deep semantic detection of generative AI phishing emails," *Expert Systems with Applications (Elsevier)*, vol. 248, p. 123491, 2025.
- [23] Y. Zhang, C. Shen, and K. Ren, "Understanding and characterizing the adoption of internationalized domain names in practice: Detection of IDN homograph attacks," *IEEE Transactions on Dependable and Secure Computing*, vol. 22, no. 1, pp. 182-196, 2024.
- [24] D. Patel, M. Joshi, and A. Shah, "Integrating cyber threat intelligence and email header forensics into gateway architectures for zero-trust phishing prevention," *IEEE Access*, vol. 14, pp. 18450-18465, 2026.
- [25] K. Al-Hubaishi, B. Mohammad, and F. Saeed, "A dual-path deep learning and header provenance framework for business email compromise (BEC) detection," *IEEE Access*, vol. 13, pp. 31205-31219, 2025.
- [26] G. Carle, F. Valenza, and M. Trevisan, "Empirical evaluation of forward-confirmed reverse DNS and cryptographic DNS-over-HTTPS in automated email authentication," in *Proc. Network Traffic Measurement and Analysis Conf. (TMA)*, 2025, pp. 78-89.
- [27] B. Klimt and Y. Yang, "The Enron corpus: A new dataset for email classification research," in *Proc. European Conf. on Machine Learning (ECML)*, 2004, pp. 217-226.
- [28] Unicode Consortium, "Unicode Technical Report #39: Unicode Security Considerations," <https://www.unicode.org/reports/tr39/>, 2024.